



# Transformers

Antoni Kwiatek  
Gradient Science Club 2025






# RNN & LSTM



# RNN




# RNN

$$a^{}=g(W_{aa}a^{}+W_{ax}x^{}+b_a)$$


$$\hat{y}^{} = g(W_{ya}a^{} + b_y)$$

$$\text{Loss}(\hat{y}^{}, y^{})$$


$$g(x) = \tanh(x)$$




$$f(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$




# RNN




# RNN




# RNN




# RNN




# RNN




# RNN




# RNN



# RNN




# RNN




$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

# RNN




$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

# RNN



$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

# RNN



$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

# RNN Problems

$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$



# RNN Problems

$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

$$\frac{1}{5} * \frac{1}{5} * \frac{1}{5} * \frac{1}{5} * \frac{1}{5} * \frac{1}{5} * \dots * \frac{1}{5} = 0.000\dots$$



# RNN Problems

$$\frac{\partial L_T}{\partial h_1} = \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_3}{\partial h_2} \cdot \dots \cdot \frac{\partial h_T}{\partial h_{T-1}} \cdot \frac{\partial L_T}{\partial h_T}$$

Vanishing Gradient

Exploding Gradient



# LSTM

LSTM with solution for Vanishing and Exploding Gradient!



# LSTM

LSTM with solution for Vanishing and Exploding Gradient!

Still not enough...



# LSTM

LSTM with solution for Vanishing and Exploding Gradient!

Still not enough...

Time complexity with backpropagation sucks

Sequential processing

Ineffective Transfer Learning


Difficult to train



# LSTM

LSTM

lient!



Provided proper attribution is provided, Google hereby grants permission to reproduce the tables and figures in this paper solely for use in journalistic or scholarly works.

---

# Attention Is All You Need

---

**Ashish Vaswani\***  
Google Brain  
avaswani@google.com

**Noam Shazeer\***  
Google Brain  
noam@google.com

**Niki Parmar\***  
Google Research  
nikip@google.com

**Jakob Uszkoreit\***  
Google Research  
usz@google.com

**Llion Jones\***  
Google Research  
llion@google.com

**Aidan N. Gomez\* †**  
University of Toronto  
aidan@cs.toronto.edu

**Lukasz Kaiser\***  
Google Brain  
lukaszkaiser@google.com

**Illia Polosukhin\* ‡**  
illia.polosukhin@gmail.com

## Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.



# Plan for Today

- RNN & LSTM roast
- *Attention is all you need*
- Deep dive into architecture
- Example workflow
- Live coding: Visualizing Attention (BERT example)
- Kahoot





# Deep dive into architecture



# Transformers architecture




Figure 1: The Transformer - model architecture.



# Transformers architecture

## Input layer





Figure 1: The Transformer - model architecture.

# Transformers architecture

Input layer

Raw text

The animal didn't cross the street because it was too tired.



# Transformers architecture


Input layer

Raw text

The animal didn't cross the street because it was too tired.

Tokenization

The animal didn't cross the street because it was too tired.



# Transformers architecture

Input layer

Raw text

The animal didn't cross the street because it was too tired.


Tokenization

The animal didn't cross the street because it was too tired.

Tokens ID (position in vocabulary)

[976, 13983, 9289, 8088, 290, 12901, 2236, 480, 673, 3101, 25920, 13]

Tokenizer - OpenAI API




# Transformers architecture

Input layer

Tokens ID (position in vocabulary)

```
[976, 13983, 9289, 8088, 290, 12901, 2236, 480, 673, 3101, 25920, 13]
```



# Transformers architecture


Input layer

Tokens ID (position in vocabulary)

```
[976, 13983, 9289, 8088, 290, 12901, 2236, 480, 673, 3101, 25920, 13]
```

321.32	321.32	321.32	321.32
1235.5	1235.5	1235.5	1235.5
2356.6	2356.6	2356.6	2356.6
0.232	0.232	0.232	0.232
.....	.....	.....	.....
123.56	123.56	123.56	123.56
3456.0	3456.0	3456.0	3456.0
45.654	45.654	45.654	45.654
1239.0	1239.0	1239.0	1239.0
54.222	54.222	54.222	54.222

Embeddings (Vector of size  $d_{model} = 512$ )





# Transformers architecture

Input layer

Positional Encoding – how the model knows where each word is


$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$



# Transformers architecture

## Encoder Layer




# Transformers architecture

Before Multi-Head Attention, there is *Self-attention*

Self-attention allows the model to relate words to each other.

## Encoder Layer

### Multi-Head Attention



# Transformers architecture


Before Multi-Head Attention, there is Self-attention

Self-attention allows the model to relate words to each other.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

## Encoder Layer

### Multi-Head Attention



# Transformers architecture

Before Multi-Head Attention, there is Self-attention


Self-attention allows the model to relate words to each other.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

*Q seeks a match in K to retrieve content from V*

## Encoder Layer


### Multi-Head Attention



# Transformers architecture

## Encoder Layer Multi-Head Attention


$$\begin{array}{ccc} \mathbf{X} & & \mathbf{W}^Q \\ \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} & \times & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \\ & & = \\ & & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \end{array}$$
$$\begin{array}{ccc} \mathbf{X} & & \mathbf{W}^K \\ \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} & \times & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \\ & & = \\ & & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \end{array}$$
$$\begin{array}{ccc} \mathbf{X} & & \mathbf{W}^V \\ \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} & \times & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \\ & & = \\ & & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \end{array}$$



# Transformers architecture


$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Softmax



## Encoder Layer

Multi-Head Attention



# Transformers architecture


$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



The	animal	didn't	cross	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
The	animal	animal	animal	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
animal	animal	animal	animal	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
didn't	didn't	didn't	didn't	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
cross	cross	cross	cross	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
the	the	the	the	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
street	street	street	street	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
because	because	because	because	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
it	it	it	it	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
was	was	was	was	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
too	too	too	too	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
tired	tired	tired	tired	the	street	because	it	was	too	tired	.
-	-	-	-	-	-	-	-	-	-	-	-
.	.	.	.	.	.	.	.	.	.	.	.

## Encoder Layer

### Multi-Head Attention




Nx



# Transformers architecture

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

The	animal	didn't	cross	the	street	because	it	was	too	tired	:		
-	0.07	0.12	0.10	0.09	0.06	0.06	0.05	0.11	0.09	0.10	0.05	0.12	
-	0.13	0.07	0.07	0.07	0.08	0.10	0.09	0.08	0.10	0.07	0.08	0.08	
didn't	0.07	0.10	0.06	0.08	0.09	0.05	0.09	0.06	0.05	0.12	0.12	0.11	
cross	-	0.07	0.06	0.11	0.08	0.06	0.09	0.06	0.13	0.07	0.10	0.07	0.09
the	-	0.08	0.06	0.12	0.10	0.12	0.11	0.08	0.12	0.05	0.06	0.05	0.06
street	-	0.07	0.07	0.11	0.07	0.07	0.09	0.06	0.11	0.05	0.13	0.11	0.06
because	0.05	0.11	0.10	0.10	0.11	0.05	0.07	0.06	0.12	0.09	0.07	0.05	
it	-	0.03	0.62	0.04	0.04	0.05	0.03	0.02	0.04	0.04	0.03	0.04	0.03
was	-	0.09	0.08	0.05	0.06	0.06	0.10	0.07	0.09	0.13	0.07	0.08	0.11
too	-	0.06	0.05	0.06	0.05	0.12	0.10	0.09	0.11	0.10	0.06	0.11	0.08
tired	-	0.11	0.12	0.07	0.06	0.06	0.08	0.11	0.12	0.05	0.08	0.08	0.06
-	-	0.05	0.07	0.12	0.07	0.08	0.10	0.07	0.12	0.12	0.06	0.08	0.06



## Encoder Layer

### Multi-Head Attention




# Transformers architecture

## Encoder Layer Multi-Head Attention

- 1) This is our input sentence\*
- 2) We embed each word\*
- 3) Split into 8 heads. We multiply  $X$  or  $R$  with weight matrices
- 4) Calculate attention using the resulting  $Q/K/V$  matrices
- 5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer




Thinking Machines

$X$



\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

$R$




# Transformers architecture



$$Output = Processed_{data} + Input$$

## Encoder Layer

Multi-Head Attention



# Transformers architecture




$$z = x + F(x)$$

$$y = \text{LayerNorm}(z)$$

## Encoder Layer

### Multi-Head Attention



# Transformers architecture

## Encoder Layer

Multi-Head Attention


Add & Norm

Residual Connection


Layer Normalization

$$z = x + F(x)$$

$$y = \text{LayerNorm}(z)$$



# Transformers architecture




$$z = x + F(x)$$

$$y = x + F(\text{LayerNorm}(x))$$

*modern models*

## Encoder Layer

### Multi-Head Attention




# Transformers architecture

## Encoder Layer Feed-Forward

### Feed-Forward


$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$




# Transformers architecture

## Encoder Layer

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$




### Encoder layer - where does it go?




# Transformers architecture

Decoder Layer



# Transformers architecture

## Decoder Layer




# Transformers architecture


Shifted right – what does it mean?

Another input and same process?

## Decoder Layer




# Transformers architecture




## Decoder Layer

Masked-Attention




# Transformers architecture



Output at a certain position can only depend  
on the words on the previous positions



## Decoder Layer

### Masked-Attention




# Transformers architecture

## Decoder Layer Masked-Attention



# Transformers architecture

Encoder layer - where does it go?



## Decoder Layer





Figure 1: The Transformer - model architecture.



# Transformers architecture


Encoder layer - where does it go?



*= Vector from Masked Multi – Head Attention*


\*Current state of output sentence

Decoder Layer



# Transformers architecture

Encoder layer - where does it go?



*$Q$  = Vector from Masked Multi – Head Attention*

\*Current state of output sentence

*$Q$  seeks a match in  $K$  to retrieve content from  $V$*

Decoder Layer




Figure 1: The Transformer - model architecture.



# Transformers architecture

## Decoder Layer





Figure 1: The Transformer - model architecture.




# Transformers architecture

Behold, a wild pi creature, foraging in its native habitat of mathematical formulas and computer code! With its infinite digits and irrational tendencies, this strange creature is beloved by mathematicians and tech enthusiasts alike. Approach with caution, for



# Decoder Layer

## Output Layer







Figure 1: The Transformer - model architecture.






2017-present









Live coding



# Questions & Discussion





Thank you!  
See you next week.

